


# Is what the desks say the authority says what the authority says?

Every one of the **533** passages the seven desks hand to an answering agent, fetched back from the publisher that wrote it and compared. **No test in this repository can answer this**, because every test reads the same stored files: a passage transcribed wrongly is stored wrongly, served wrongly and asserted wrongly, and every green stays green.



The same passage, two roads. They meet at the comparison, and the verdict is read off that block rather than written ahead of it.

## The roster

|                                 |           |            |           |                      |                   |
|---------------------------------|-----------|------------|-----------|----------------------|-------------------|
| 2                               | 0         | 531 of 533 |           | 13                   |                   |
| DIFFERS                         | COULD NOT | TIED       |           | TIED ON SPACING ONLY |                   |
| DESK                            | PASSAGES  | DIFFERS    | COULD NOT | TIED                 | OF THOSE, SPACING |
| capitalization-and-de-minimis   | 58        | 1          |           | 57                   |                   |
| cash-and-bank                   | 49        | 1          |           | 48                   | 4                 |
| fixed-assets                    | 172       |            |           | 172                  |                   |
| meals-and-entertainment         | 44        |            |           | 44                   |                   |
| personal-or-business            | 34        |            |           | 34                   | 1                 |
| rewards-and-information-returns | 47        |            |           | 47                   | 8                 |
| vehicle-expense                 | 129       |            |           | 129                  |                   |

Read the two differences first; the 531 that agree are not what anybody needs to read. Each desk has its own exhibit beside this one, with its sample of one worked link by link.

## What it found

Two differences in 533 passages, and a defect in the record that no difference would have shown.

- **Right text, wrong link — three passages.** The rewards desk stores § 6041, § 6041A, § 6050W and § 6071 under one source whose recorded URL is the House's granule for **§ 6041 only**. A reader following that link finds the first and not the other three; the "6041A" and "6050W" they see on that page are § 6041's own cross-references. All three bodies are **verbatim correct** on their own granules, so nothing served was wrong and no test could have caught it — the numbers were right and the citation pointed somewhere that does not carry them. **Only following the link finds this.** The checker now resolves each statute citation to its own granule; the record's source URL is still under-covering and is a fix for the desk.
- **A partial quotation that does not say it is one.** Publication 583's sentence has two branches and the cash desk stores the second without marking the omission. The split is right — the branches have opposite answers, which is the defect #264 found — but an answerer reading the passage cannot tell a branch was removed.
- **A publisher's interface text inside a quotation.** irs.gov appends the word **PDF** to links that point at PDFs, and it sits mid-sentence on the tangible-property page. Ours drops it, which reads correctly and is not what the page literally shows.

**And what it did not find is the point of running it.** 531 of 533 passages are, character for character, what the publisher publishes today. That was an assumption this morning and it is a measurement now.

---

## What I got wrong

**The first run reported 88 differences and 34 of them were mine.** The extractor replaced every markup tag with a space, so eCFR's election</I>—(1) came back as election - (1) and 601.601(d)(2)(ii)(<I>b</I>) as ( b ). A checker that invents differences is worse than no checker, because somebody goes and "fixes" the record to match it. Inline tags now close up and only block tags become a space.

**Then it reported 18, and 16 of those were mine too.** pypdf reads a kerned two-column IRS page as Y ou and hyphenates infor-mation across a line break, and irs.gov puts block markup inside a sentence. Those are rendering artefacts of the extractor, not differences in what was published — so the comparison now runs twice, and the exhibit says which pass carried each line rather than quietly folding them together.

**I could not capture the sources as photographs, and the skill asks for photographs.**

This machine's browser cannot reach the publishers even when pointed at the egress proxy — ERR\_CONNECTION\_RESET — while the HTTP client reaches them fine. Attempted, failed, recorded. What is here instead is the document itself: its SHA-256, its length, the moment it was fetched and the publisher's own words quoted around each match. For text that is the stronger evidence, because a reader can recompute the hash and a screenshot cannot be re-verified at all.

---

## What this does not prove

- **It proves the words, not the reading.** That § 1.263(a)-1(f)(1)(ii)(D) says \$500 is now established. Whether \$500 is the number that governs a client is a question about Notice 2015-82 and about the firm's election, and no fetch settles it.

- **It proves containment, not completeness.** The check asks whether our passage occurs in the publisher's document. It does not ask whether we stored the *whole* of what that citation covers — a passage could be a true fragment of a paragraph whose second half changes its meaning. Publication 583 above is exactly that failure caught by accident, not by design.
- **The whitespace-ignoring pass can hide a genuinely missing space.** 13 lines tied only on that pass. Each is named in the roster as spacing so nobody has to take it on trust.
- **The regulations were fetched as of 2026-01-01, not today.** eCFR's versioner refuses a future date and this is the last date known good on it. A regulation amended between that date and now would not show here.
- **One reading, not two.** The skill asks for a second rendering where one exists — the XML and the rendered page, the PDF and the HTML. Only one rendering per source was read. A digit or a word wrong in *both* our copy and the one rendering we read would pass.
- **Nothing here touches whether a desk answers well.** This is about the authority it is given. A desk with a perfect corpus can still reason badly from it, and that is what the scoreboard is for.