


Is what the desks say the authority says what the authority says?

Every one of the **705** passages the seven desks hand to an answering agent, fetched back from the publisher that wrote it and compared. **No test in this repository can answer this**, because every test reads the same stored files: a passage transcribed wrongly is stored wrongly, served wrongly and asserted wrongly, and every green stays green.



The same passage, two roads. They meet at the comparison, and the verdict is read off that block rather than written ahead of it.

The roster

0	0	705 of 705		13	
DIFFERS	COULD NOT	TIED	TIED ON SPACING ONLY		
DESK	PASSAGES	DIFFERS	COULD NOT	TIED	OF THOSE, SPACING
capitalization-and-de-minimis	89			89	
cash-and-bank	49			49	4
fixed-assets	289			289	
meals-and-entertainment	44			44	
personal-or-business	34			34	1
rewards-and-information-returns	61			61	8
vehicle-expense	139			139	

Nothing here differs, so read what it took to get there rather than the roster. The 705 that agree are not what anybody needs to read. Each desk has its own exhibit beside this one, with its sample of one worked link by link.

What it found

Nothing, and that sentence had to be earned three times over. All 531 passages the seven desks hand an answering agent are, today, what the publisher publishes: **0 differs, 0 could not**. The run this morning said 533 and two; the one before it said 88 differences. What changed in between is the record and the checker, and both are worth reading before this figure is believed.

The three findings this run closed, each fixed in the record rather than in the check.

- **Right text, wrong link — three passages.** The rewards desk stored § 6041, § 6041A, § 6050W and § 6071 under ONE source whose recorded URL is the House's granule for **§ 6041 only**. All four bodies were **verbatim correct**, so nothing served was ever wrong and no test in the repository could have seen it — a reader following the citation for three of the four simply landed on a page that did not contain it. **Only opening the link finds this**. Split into four sources, four URLs, each verified live; and the record now has a test comparing the URL a citation is really on with the URL its source claims.
- **A partial quotation that did not say it was one.** Publication 583's sentence has two branches and the cash desk stored the second without marking the omission. The split is right — the branches have opposite answers, which is the defect #264 found — but an answerer could not tell a branch had been removed. Marked now with [. . .], and the mark is CHECKED rather than excused: every segment must appear in the live document, in order.
- **A publisher's interface text inside a quotation.** irs.gov appends the word **PDF** to links that point at PDFs, and one sits mid-sentence on the tangible-property page. Ours dropped it. This was judged in our favour this morning and left; it is the same unmarked omission as Pub 583, and it is marked now. It was the last difference in the corpus.

What "nothing" is worth here. Three of the 531 tied only because a marked omission was checked segment by segment, and 13 tied only once whitespace was ignored. Both are declared, both are counted on each desk's own sheet, and neither is silently folded into the headline. A clean run that hides how it got clean is the thing this exhibit exists not to be.

What I got wrong

The first run reported 88 differences and 34 of them were mine. The extractor replaced every markup tag with a space, so eCFR's election</I>—(1) came back as election - (1) and 601.601(d)(2)(ii)(<I>b</I>) as (b). A checker that invents differences is worse than no checker, because somebody goes and "fixes" the record to match it. Inline tags now close up and only block tags become a space.

Then it reported 18, and 16 of those were mine too. pypdf reads a kerned two-column IRS page as Y ou and hyphenates infor-mation across a line break, and irs.gov puts block markup inside a sentence. Those are rendering artefacts of the extractor, not differences in what was published — so the comparison now runs twice, and the exhibit says which pass carried each line rather than quietly folding them together.

I could not capture the sources as photographs, and the skill asks for photographs.

This machine's browser cannot reach the publishers even when pointed at the egress proxy — `ERR_CONNECTION_RESET` — while the HTTP client reaches them fine. Attempted, failed, recorded. What is here instead is the document itself: its SHA-256, its length, the moment it was fetched and the publisher's own words quoted around each match. For text that is the stronger evidence, because a reader can recompute the hash and a screenshot cannot be re-verified at all.

What this does not prove

- **It proves the words, not the reading.** That § 1.263(a)-1(f)(1)(ii)(D) says \$500 is now established. Whether \$500 is the number that governs a client is a question about Notice 2015-82 and about the firm's election, and no fetch settles it.
- **It proves containment, not completeness.** The check asks whether our passage occurs in the publisher's document. It does not ask whether we stored the *whole* of what that citation covers — a passage could be a true fragment of a paragraph whose second half changes its meaning. Publication 583 above is exactly that failure caught by accident, not by design.
- **The whitespace-ignoring pass can hide a genuinely missing space.** 13 lines tied only on that pass. Each is named in the roster as spacing so nobody has to take it on trust.
- **The regulations were fetched as of 2026-08-28, not today.** eCFR's versioner refuses a future date and this is the last date known good on it. A regulation amended between that date and now would not show here.
- **One reading, not two.** The skill asks for a second rendering where one exists — the XML and the rendered page, the PDF and the HTML. Only one rendering per source was read. A digit or a word wrong in *both* our copy and the one rendering we read would pass.
- **Nothing here touches whether a desk answers well.** This is about the authority it is given. A desk with a perfect corpus can still reason badly from it, and that is what the scoreboard is for.